

Fraud detector model training

```
In [1]: import sys
import types
import pickle

import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import sklearn.metrics as metrics
from sklearn.metrics import roc_auc_score, confusion_matrix, precision_recall_curve
from sklearn.metrics import confusion_matrix, f1_score
import matplotlib.pyplot as plt
```

```
In [2]: transactions = pd.read_csv('creditcard.csv')
```

```
In [3]: transactions.head()
```

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.36
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.25
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.57
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.38
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.87

5 rows × 31 columns

```
In [4]: number_of_rows = transactions.shape[0]
transactions.shape
```

Out[4]: (284807, 31)

```
In [5]: fraud_cases = transactions[(transactions.Class==1)]
        fraud_cases.head()
```

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8
541	406.0	-2.312227	1.951992	-1.609851	3.997906	-0.522188	-1.426545	-2.537387	1.391657
623	472.0	-3.043541	-3.157307	1.088463	2.288644	1.359805	-1.064823	0.325574	-0.067794
4920	4462.0	-2.303350	1.759247	-0.359745	2.330243	-0.821628	-0.075788	0.562320	-0.399147
6108	6986.0	-4.397974	1.358367	-2.592844	2.679787	-1.128131	-1.706536	-3.496197	-0.248778
6329	7519.0	1.234235	3.019740	-4.304597	4.732795	3.624201	-1.357746	1.713445	-0.496358

5 rows × 31 columns

```
In [6]: number_of_fraud_cases = fraud_cases.shape[0]
        fraud_cases.shape
```

Out[6]: (492, 31)

Important Note

The classes are highly imbalanced!

```
In [7]: percentage_of_fraud_transactions = number_of_fraud_cases / number_of_rows *
        print('Percentage of fraud transactions', percentage_of_fraud_transactions,
              Percentage of fraud transactions 0.1727485630620034 %
```

```
In [8]: X = transactions.drop('Class', 1)
        y = transactions['Class']
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
```

Simple random forest classifier

```
In [9]: %%time
rf_model = RandomForestClassifier(
    n_estimators=1,
    criterion='gini',
    max_depth=7,
    min_samples_split=2,
    min_samples_leaf=5,
    min_weight_fraction_leaf=0.0,
    max_features='auto',
    max_leaf_nodes=None,
    bootstrap=True,
    oob_score=False,
    n_jobs=16,
    random_state=None,
    verbose=100,
    warm_start=False,
    class_weight=None).fit(X_train, y_train)
```

```
[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent
workers.
building tree 1 of 1
[Parallel(n_jobs=16)]: Done    1 tasks      | elapsed:    0.6s
[Parallel(n_jobs=16)]: Done    1 out of   1 | elapsed:    0.6s finished
CPU times: user 624 ms, sys: 15.8 ms, total: 640 ms
Wall time: 656 ms
```

```
In [10]: pred_train = rf_model.predict(X_train)
pred_test = rf_model.predict(X_test)
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent w
orkers.
[Parallel(n_jobs=1)]: Done    1 out of   1 | elapsed:    0.0s remaining:
0.0s
[Parallel(n_jobs=1)]: Done    1 out of   1 | elapsed:    0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent w
orkers.
[Parallel(n_jobs=1)]: Done    1 out of   1 | elapsed:    0.0s remaining:
0.0s
[Parallel(n_jobs=1)]: Done    1 out of   1 | elapsed:    0.0s finished
```

```
In [11]: confusion_matrix(y_train, pred_train)
```

```
Out[11]: array([[227423,    30],
               [    91,   301]])
```

```
In [12]: precision, recall, _ = precision_recall_curve(y_test, pred_test)
print('Precision: ', precision[1])
print('Recall: ', recall[1])
print('AUC', metrics.auc(precision, recall))
```

```
Precision:  0.9125
Recall:  0.73
AUC 0.8197314437695304
```

Improve random forest model using grid search

```
In [13]: rf_model = RandomForestClassifier(
    n_estimators=1,
    criterion='gini',
    max_depth=None,
    min_samples_split=2,
    min_samples_leaf=5,
    min_weight_fraction_leaf=0.0,
    max_features='auto',
    max_leaf_nodes=None,
    bootstrap=True,
    oob_score=False,
    n_jobs=16,
    random_state=None,
    verbose=0,
    warm_start=False,
    class_weight=None)
```

```
In [14]: param_grid = {
    'n_estimators': [6, 7, 8, 9, 10, 20, 30, 50],
    'max_depth': [7, 8, 9, 10, 11]
}
```

```
In [15]: clf = GridSearchCV(
    rf_model,
    param_grid,
    n_jobs=16,
    cv=3,
    scoring='recall'
)
```

```
In [16]: %%time
         clf.fit(X_train, y_train)
```

CPU times: user 10.1 s, sys: 445 ms, total: 10.5 s
Wall time: 4min 14s

```
Out[16]: GridSearchCV(cv=3, error_score='raise-deprecating',
                    estimator=RandomForestClassifier(bootstrap=True, class_weigh
t=None,
                                                    criterion='gini', max_depth
=None,
                                                    max_features='auto',
                                                    max_leaf_nodes=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=5,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.
0,
                                                    n_estimators=1, n_jobs=16,
                                                    oob_score=False,
                                                    random_state=None, verbose=
0,
                                                    warm_start=False),
                    iid='warn', n_jobs=16,
                    param_grid={'max_depth': [7, 8, 9, 10, 11],
                                'n_estimators': [6, 7, 8, 9, 10, 20, 30, 50]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                    scoring='recall', verbose=0)
```

```
In [17]: rf_best = clf.best_estimator_
         pred_train = rf_best.predict(X_train)
         pred_test = rf_best.predict(X_test)
         confusion_matrix(y_train, pred_train)
```

```
Out[17]: array([[227441,    12],
               [    80,   312]])
```

```
In [18]: confusion_matrix(y_test, pred_test)
```

```
Out[18]: array([[56857,    5],
               [   22,   78]])
```

```
In [19]: precision, recall, _ = precision_recall_curve(y_train, pred_train)
         print('Precision: ', precision[1])
         print('Recall: ', recall[1])
```

Precision: 0.9629629629629629
Recall: 0.7959183673469388

```
In [20]: precision, recall, _ = precision_recall_curve(y_test, pred_test)
print('Precision: ', precision[1])
print('Recall: ', recall[1])
```

Precision: 0.9397590361445783
Recall: 0.78

```
In [21]: print('AUC', metrics.auc(precision, recall))
print('The general score for rf model on a test set:', rf_best.score(X_test, y_test))
print('F1 score for the rf model on a test set: ', f1_score(y_test, pred_test))
```

AUC 0.8583170729334247
The general score for rf model on a test set: 0.9995259997893332
F1 score for the rf model on a test set: 0.8524590163934427

Let's try gradient boosting

```
In [22]: from catboost import CatBoostClassifier
from catboost import Pool
from catboost import MetricVisualizer
```

```
In [23]: # create train_pool object
train_pool = Pool(
    data=X_train,
    label=y_train
)

# create validation_pool object
validation_pool = Pool(
    data=X_test,
    label=y_test
)
```

```
In [24]: # we create the object of CatBoostClassifier class
cbs = CatBoostClassifier(iterations=4000,
                        learning_rate=0.008,
                        custom_loss=['AUC', 'Accuracy', 'Recall'],
                        max_depth=10)

# we are passing categorical features as parameters here
# verbose = 10 outputs only each 10th tree

cbs.fit(
    train_pool,
    eval_set=validation_pool,
    verbose=False,
    plot=True
);
```

☒ --- Learn ☒ — Eval


Logloss AUC Accuracy Recall

☒ catboost_info 1m 48s

--- learn — test

curr --- 0.8571428... — 0.8 784

best 0.82 1476



☐ Click Mode ☐ Logarithm

☐ Smooth

```
In [25]: pred_test_cbs = cbs.predict(X_test)
confusion_matrix(y_test, pred_test_cbs)
```

```
Out[25]: array([[56858,    4],
               [   18,   82]])
```

```
In [26]: precision_cbs, recall_cbs, _ = precision_recall_curve(y_test, pred_test_cbs)
print('AUC', metrics.auc(precision_cbs, recall_cbs))
print('Precision: ', precision_cbs[1])
print('Recall: ', recall_cbs[1])
print('The accuracy for cbs model on a test set:', cbs.score(X_test, y_test))
print('F1 score for the cbs model on a test set: ', f1_score(y_test, pred_t
```

```
AUC 0.885146629780931
Precision:  0.9534883720930233
Recall:  0.82
The accuracy for cbs model on a test set: 0.9996137776061234
F1 score for the cbs model on a test set:  0.8817204301075269
```

Note: the best result for max_depth=6 model was AUC 0.879. So increasing the depth of the trees actually improves the classifier.

Note: in the models above we have dealt with unbalanced dataset in a pretty dummy way, so we really have to revisit the model creation

Side note: Comprehensive Kernel on unbalanced datasets oversampling.

<https://www.kaggle.com/janiobachmann/credit-fraud-dealing-with-imbalanced-datasets>
<https://www.kaggle.com/janiobachmann/credit-fraud-dealing-with-imbalanced-datasets>

Save both models

First, will use ordinary pickling to serialize the models.

```
In [27]: filename = 'random_forest_model.sav'
pickle.dump(rf_best, open(filename, 'wb'))
```

```
In [28]: filename = 'catboost_model.sav'
pickle.dump(cbs, open(filename, 'wb'))
```

```
In [29]: import os
def print_file_size(filename):
    statinfo = os.stat(filename)
    print(filename, ': ', statinfo.st_size/1000, 'Kb')

print_file_size('random_forest_model.sav')
print_file_size('catboost_model.sav')
```

```
random_forest_model.sav : 55.277 Kb
catboost_model.sav : 25664.676 Kb
```

Type *Markdown* and LaTeX: α^2

